

Cross-Validation in Practice (Summary)

Biostatistics 735/794

Standard Validation (Single Split)

- ▶ Split data once into three parts:
 - ▶ **Training set** — used to fit models.
 - ▶ **Validation set** — used to tune hyperparameters and choose the best model.
 - ▶ **Test set** — held out until the very end to estimate generalization error.
- ▶ **Pros:**
 - ▶ Simple to implement and fast.
 - ▶ Mimics the “future data” scenario with a true test set.
- ▶ **Cons:**
 - ▶ High variability — results depend on a single random split.
 - ▶ Wastes data, since the test set never contributes to training.
- ▶ **Use case:** Large datasets where a single split still leaves plenty of training data.

1. Standard Cross-Validation

- ▶ Split data into K folds.
- ▶ Rotate: each fold is validation once, others are training.
- ▶ **Uses:**
 - ▶ Compare multiple candidate models (choose best).
 - ▶ Estimate error of a single fixed model.
- ▶ **Key idea:** Every observation is used for both training and validation across iterations.

2. CV + Test Dataset

- ▶ Hold out a test set from the start.
- ▶ Run CV only on training set to choose/tune model.
- ▶ After choosing, refit the final model on the entire training set
- ▶ Predict and evaluate the error of the test set.
- ▶ **Pros:**
 - ▶ Honest estimate of generalization error.
 - ▶ Avoids reusing the same data for both tuning and final error.
- ▶ **Cons:**
 - ▶ Test set must be large enough to be reliable.
 - ▶ Less data available for training.

3. Nested Cross-Validation

- ▶ Outer loop: split into K_{outer} folds.
- ▶ For each outer training fold: run an **inner CV** to choose/tune the model.
- ▶ Evaluate tuned model on the outer validation fold.
- ▶ Average outer validation errors for the overall estimate.
- ▶ **Pros:**
 - ▶ Uses all data for both tuning and error estimation.
 - ▶ Honest error estimate (no leakage).
- ▶ **Cons:**
 - ▶ Computationally expensive.
 - ▶ Harder to explain and implement.

4. Stratified Cross-Validation

- ▶ Ensure folds preserve outcome distribution.
 - ▶ For classification: same class proportions per fold.
 - ▶ For continuous outcomes: bin into quantiles, then stratify.
- ▶ **Use:** Reduce variability of error estimates, especially with imbalanced outcomes.
- ▶ **Note:** This approach can be used to estimate the error of a single model or combined with *CV + Test Dataset* or *Nested CV*.

5. Group-Specific Error Estimation

- ▶ After CV, compute error separately within subgroups.
 - ▶ e.g., Age bands, sex, clinical strata.
- ▶ **Use:** Check fairness, detect whether the model works equally well across groups.
- ▶ **Note:** This approach can be used to estimate the error of a single model or combined with *CV + Test Dataset* or *Nested CV*.

Key Takeaways

- ▶ Always separate the roles of training, validation, and testing.
- ▶ Standard CV is fine for model selection or error estimation.
- ▶ If you want an honest estimate of prediction error *after model selection*, use either:
 - ▶ CV + independent test set, or
 - ▶ Nested CV.
- ▶ Use stratification to stabilize results, and report group-specific errors to evaluate fairness.